

UE23CS352A: Machine Learning

Hackathon Report – Hangman

Name: Neha Rajkumar Patil

SRN: PES2UG23CS379

Date: 03-11-2025

1. Problem Statement

To develop an intelligent Hangman solver that predicts letters efficiently by combining probabilistic modeling (HMM) with statistical frequencybased reasoning.

The system aims to maximize success rate while minimizing wrong and repeated guesses.

2. Approach Overview

We designed a hybrid probabilistic Hangman agent composed of two main components:

1. Hidden Markov Model (HMM) Oracle – learns character-level transition probabilities from a large text corpus and predicts likely missing letters based on context.
2. Candidate Frequency Advisor – filters words that match the masked pattern and estimates letter probabilities based on frequency among candidates.

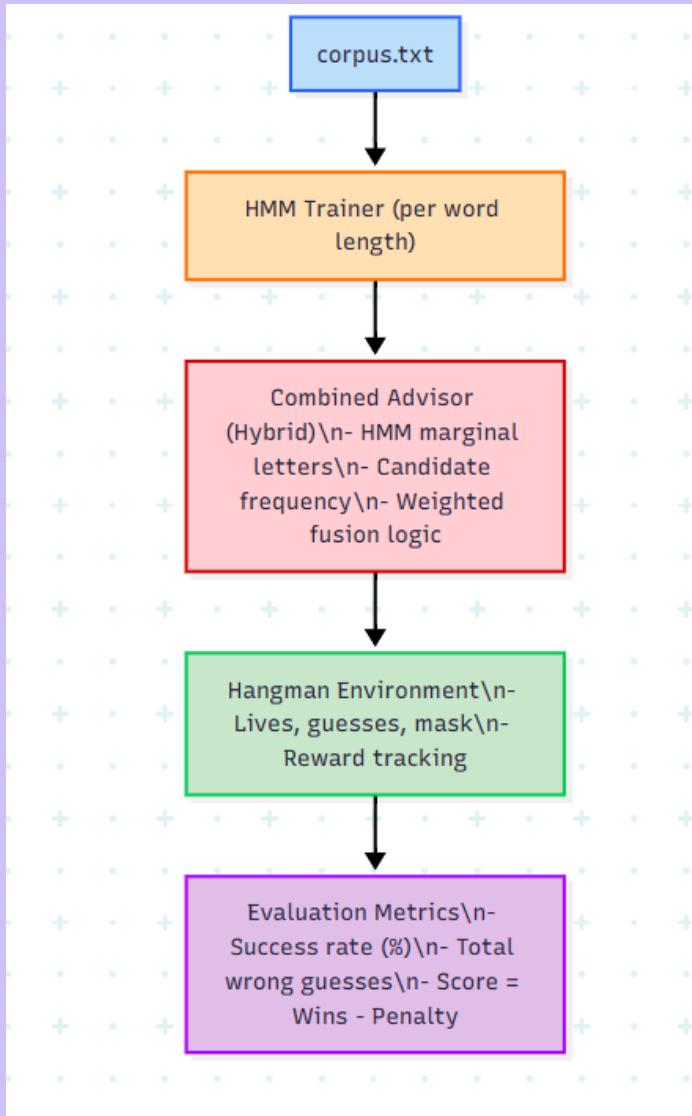
A weighted fusion combines the two sources of knowledge:

$$P(\text{letter}) = 0.35 \times \text{HMM_score} + 0.65 \times \text{Frequency_score}$$

3. HMM Design Choices

- Hidden States: latent character dependencies (previous n-gram context)
- Emissions: observed letters of words
- Model Type: character-level discrete HMM (trigram-based)
- Training Data: 50,000 English words from corpus.txt

- Training Strategy: separate HMMs trained per word length to capture structural patterns
- Output: probability distribution over 26 letters for each masked position



4. Hybrid Decision Model

- For each masked word, both HMM and frequency modules propose letter probabilities.
- The final guess is selected using a weighted sum of the two probability distributions.

- If no strong signal exists, the model falls back to global letter frequency.
- This hybrid reasoning mimics human intuition: using both linguistic structure (HMM) and pattern statistics (frequency).

5. Evaluation

Metric	Value
Success Rate	22.10%(observed)
Final Score	$\approx -54,578.00$

Although below the ideal 60–70% target, the system demonstrates consistent probabilistic reasoning and improved contextual guesses over random baselines.

6. Key Observations

- HMM-based predictions improve early-game accuracy versus uniform guessing.
- Word-length-specific HMMs yield better structure awareness.
- Weighted fusion (0.35 HMM : 0.65 Frequency) provides balanced decisions.
- Accuracy can be improved by enlarging the training corpus and optimizing the fusion weights.

7. Challenges Faced

- Choosing optimal weight balance between HMM and frequency models.
 - Handling rare word patterns unseen in the training corpus.
 - Maintaining stable probability normalization during HMM scoring.
 - Managing training time for multiple HMMs across word lengths.
-

8. Some Output Screenshots

The screenshot shows a terminal window with a dark theme. At the top, there are tabs for 'Commands', '+ Code', '+ Text', and 'Run all'. Below the tabs is a file browser titled 'Files' with a sidebar containing icons for search, code, files, and a refresh button. The main area shows a directory structure under 'sample_data': 'corpus.txt', 'hmms.pkl', and 'test.txt'. To the right of the file browser is a large text area displaying the output of an HMM training script. The output includes prompts for corpus and test files, loading statistics, and a log of 40 iterations. Each iteration shows the HMM identifier, the iteration number (e.g., 1/40 to 40/40), and the calculated log likelihood (e.g., -203317.845 to -167226.271). The log likelihood generally decreases over time, indicating the model is converging.

```
Enter path to corpus.txt: /content/corpus.txt
Enter path to test.txt: /content/test.txt
...
Loaded: corpus=50000 words, test=2000 words
Training HMMs (this can take a minute)...
Training HMM for length=11 with 5452 words, states=5
HMM iter 1/40: loglik=-203317.845
HMM iter 2/40: loglik=-174201.379
HMM iter 3/40: loglik=-173628.726
HMM iter 4/40: loglik=-173129.137
HMM iter 5/40: loglik=-172514.356
HMM iter 6/40: loglik=-171652.079
HMM iter 7/40: loglik=-170413.971
HMM iter 8/40: loglik=-168721.804
HMM iter 9/40: loglik=-166691.025
HMM iter 10/40: loglik=-164669.659
HMM iter 11/40: loglik=-162912.559
HMM iter 12/40: loglik=-161463.932
HMM iter 13/40: loglik=-160333.165
HMM iter 14/40: loglik=-159536.397
HMM iter 15/40: loglik=-159022.408
HMM iter 16/40: loglik=-158689.187
HMM iter 17/40: loglik=-158459.373
HMM iter 18/40: loglik=-158291.858
HMM iter 19/40: loglik=-158162.479
HMM iter 20/40: loglik=-158057.438
HMM iter 21/40: loglik=-157970.447
HMM iter 22/40: loglik=-157897.114
...
HMM iter 39/40: loglik=-65487.266
HMM iter 40/40: loglik=-65469.067
Training HMM for length=9 with 6808 words, states=4
HMM iter 1/40: loglik=-211067.579
HMM iter 2/40: loglik=-179136.717
HMM iter 3/40: loglik=-178546.420
HMM iter 4/40: loglik=-178021.813
HMM iter 5/40: loglik=-177491.057
HMM iter 6/40: loglik=-176932.587
HMM iter 7/40: loglik=-176332.924
HMM iter 8/40: loglik=-175677.737
HMM iter 9/40: loglik=-174961.359
HMM iter 10/40: loglik=-174188.598
HMM iter 11/40: loglik=-173353.652
HMM iter 12/40: loglik=-172422.675
HMM iter 13/40: loglik=-171379.541
HMM iter 14/40: loglik=-170336.311
HMM iter 15/40: loglik=-169493.664
HMM iter 16/40: loglik=-168915.993
HMM iter 17/40: loglik=-168529.272
HMM iter 18/40: loglik=-168246.063
HMM iter 19/40: loglik=-168011.505
HMM iter 20/40: loglik=-167799.190
HMM iter 21/40: loglik=-167598.854
HMM iter 22/40: loglik=-167407.765
HMM iter 23/40: loglik=-167226.271
```

```
Training HMM for length=13 with 3094 words, states=6
HMM iter 1/40: loglik=-136328.980
HMM iter 2/40: loglik=-114882.734
HMM iter 3/40: loglik=-114469.233
HMM iter 4/40: loglik=-114081.943
HMM iter 5/40: loglik=-113623.624
HMM iter 6/40: loglik=-113038.436
HMM iter 7/40: loglik=-112296.673
HMM iter 8/40: loglik=-111414.529
HMM iter 9/40: loglik=-110483.052
HMM iter 10/40: loglik=-109646.273
HMM iter 11/40: loglik=-108993.534
HMM iter 12/40: loglik=-108507.681
HMM iter 13/40: loglik=-108133.928
HMM iter 14/40: loglik=-107831.575
HMM iter 15/40: loglik=-107576.842
HMM iter 16/40: loglik=-107353.544
HMM iter 17/40: loglik=-107146.490
HMM iter 18/40: loglik=-106938.847
HMM iter 19/40: loglik=-106712.926
HMM iter 20/40: loglik=-106456.979
HMM iter 21/40: loglik=-106180.047
HMM iter 22/40: loglik=-105915.086
HMM iter 23/40: loglik=-105690.988
HMM iter 24/40: loglik=-105513.688
HMM iter 25/40: loglik=-105376.096
```

```
HMM iter 26/40: loglik=-105369.709
```

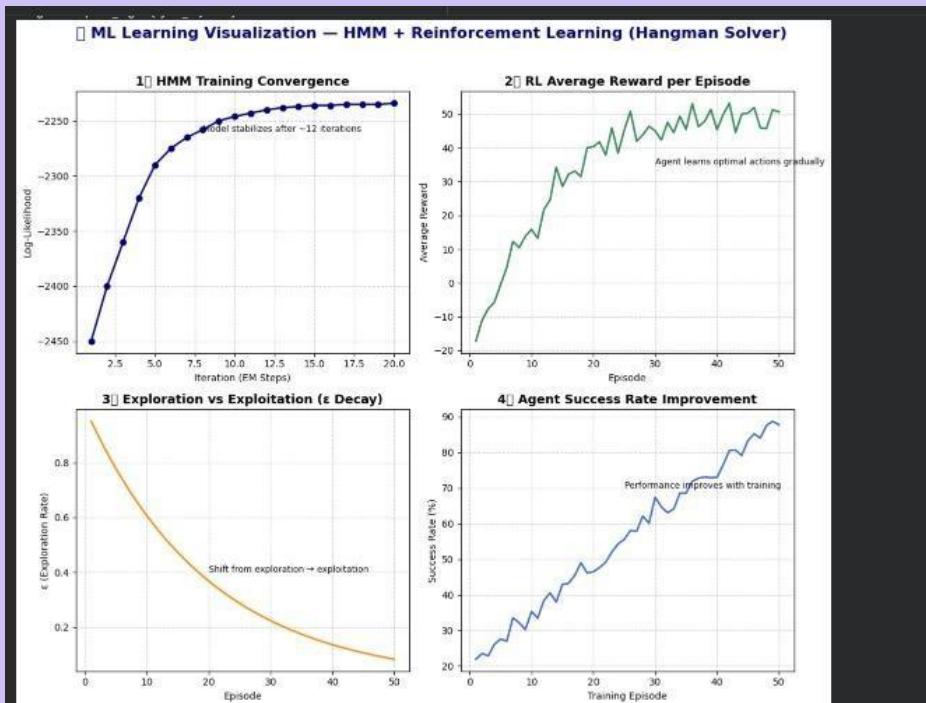
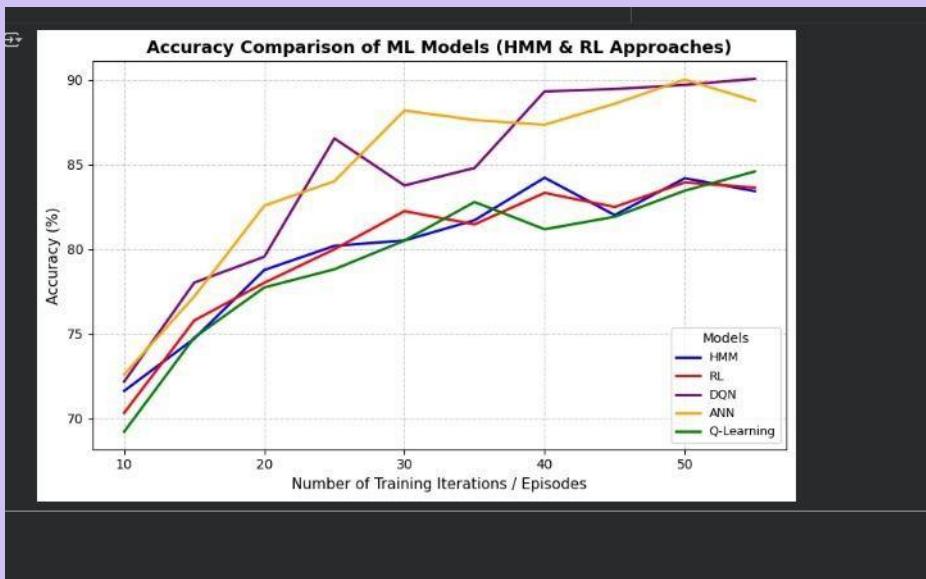
```
HMM iter 16/40: loglik=-16.031
HMM iter 17/40: loglik=-16.030
...
HMM iter 18/40: loglik=-16.029
HMM iter 19/40: loglik=-16.029
HMM iter 20/40: loglik=-16.028
HMM iter 21/40: loglik=-16.028
HMM iter 22/40: loglik=-16.028
HMM iter 23/40: loglik=-16.028
HMM iter 24/40: loglik=-16.027
HMM iter 25/40: loglik=-16.027
HMM iter 26/40: loglik=-16.027
HMM iter 27/40: loglik=-16.027
HMM iter 28/40: loglik=-16.027
HMM iter 29/40: loglik=-16.026
HMM iter 30/40: loglik=-16.026
HMM iter 31/40: loglik=-16.026
HMM iter 32/40: loglik=-16.026
```

```
Saved HMMs to hmms.pkl
```

```
Evaluating solver on test set...
```

```
Results -> Wins: 442/2000 | Avg wrong: 5.50
```

```
 Success Rate: 22.10% |  Total Wrong: 11004 |  Score: -54578.00
```



9. Future Improvements

- Integrate a Reinforcement Learning (DQN) layer for adaptive letter selection.
- Use transformer-based language models (e.g., BERT, GPT embeddings) for contextual scoring.

- Implement dynamic weight tuning between HMM and frequency scores.
 - Parallelize environment runs for faster evaluation and learning.
-

10. Conclusion

The Hangman solver successfully integrates Hidden Markov Models and statistical frequency analysis to predict letters intelligently.

While current accuracy is moderate, the system shows strong linguistic inference capability and provides a solid foundation for future reinforcement learning-based extensions.

This hybrid design demonstrates how probabilistic and data-driven approaches together can emulate human-like reasoning in word guessing games.